

C++ Pointers (Part I)



Krishna Kumar

Every computer, at the unreachable memory address 0x-1, stores a secret. I found it, and it is that all humans ar-- SEGMENTATION FAULT
XKCD (<https://xkcd.com/138/>)

Pointer

- Pointer values are memory addresses
 - A pointer **p** can hold the address of a memory location
 - A pointer points to an object of a given type
 - E.g. a ***int**** points to an ***int***, not to a ***string***

Pointer to

<i>char* cptr;</i>	// Pointer to a char
<i>int* iptr;</i>	// Pointer to an int
<i>float* fptr;</i>	// Pointer to a float
MyClass* myclasspt;	// Pointer to a user-defined class MyClass
<i>int* ap[15];</i>	// array of 15 pointers to ints
<i>int (*fp)(char*);</i>	// pointer to function taking a char* // argument; returns an int
<i>int* f(char*);</i>	// function taking an int* argument; returns a // pointer to int

Pointer Danger

int ptr;* // create a pointer-to-int

**ptr = 556;* // place a value in never-never land

- **Pointer Golden Rule:** Always initialize a pointer to a definite and appropriate address before you apply the dereferencing operator (*) to it

Capsule Summary

<i>int v;</i>	//defines variable v of type int
<i>int* p = &v;</i>	//defines p as a pointer to int
	//assigns address of variable v
to	
	// pointer p
<i>v = 3;</i>	//assigns 3 to v
<i>*p = 3;</i>	//also assigns 3 to v

Access

- A pointer does **not** know the number of elements that it's pointing to (only the address of the first element)


```
double* p1 = new double;
```

```
*p1 = 7.3;           // ok
```

```
p1[0] = 8.2;        // ok
```

```
p1[17] = 9.4;       // ouch! Undetected error
```

```
p1[-4] = 2.4;       // ouch! Another undetected error
```



```
double* p2 = new double[100]; // !Warning: Avoid magic numbers like 100
```

```
// when you write real code (BAD)
```

```
*p2 = 7.3;           // ok
```

```
p2[17] = 9.4;        // ok
```

```
p2[-4] = 2.4;       // ouch! Undetected error
```




Access

- A pointer does **not** know the number of elements that it's pointing to


```
double* p1 = new double;
```

```
double* p2 = new double[100];
```



```
p1[17] = 9.4;    // error (obviously)
```

```
p1 = p2;        // assign the value of p2 to p1
```



```
p1[17] = 9.4;    // now ok: p1 now points to the array of 100 doubles
```

Access

- A pointer **does** know the type of the object that it's pointing to

int pi1 = new int(7);*

int pi2 = pi1;* *// ok: pi2 points to the same object as pi1*

double pd = pi1;* *// error: can't assign an int* to a double**

char pc = pi1;* *// error: can't assign an int* to a char**

- There are no implicit conversions between a pointer to one value type to a pointer to another value type
- However, there are implicit conversions between value types:

**pc = 8;* *// ok: we can assign an int to a char*

**pc = *pi1;* *// ok: we can assign an int to a char*

*int * pi;*

pi = 0xB8000000; *// type mismatch (You know that's an address compiler doesn't)*

pi = (int) 0xB8000000;* *// types now match (int 2 Bytes and addr may be 4Bytes)*

Computer Memory

- **Stack:**
 - Stores local data, return addresses, used for parameter passing
- **Heap:**
 - You would use the heap if you don't know exactly how much data you will need at runtime or if you need to allocate a lot of data
- **Use local variables (stack) when you can. Use dynamic allocation (heap) when you have to.**



Stack Memory

- Variables when out of scope will **automatically be deallocated.**
- **Faster to allocate** in comparison to variables on the heap.
- Can have a **stack overflow** when too much of the stack is used.
- Data created on the stack **can be used without pointers.**
- Use stack when data is not too big.

Heap Memory

- Variables on the heap must be **destroyed manually** and **never fall out of scope**.
- The data is freed with **delete, delete[]** or **free**
- **Slower to allocate** in comparison to variables on the stack.
- Data can be accessed by pointers
- Can have allocation failures if too big of a buffer is requested to be allocated.
- You would use the heap if you don't know exactly how much data you will need at runtime or if you need to allocate a lot of data.
- **Responsible for memory leaks**

The free store

(sometimes called "the heap")

- You request memory "to be allocated" "on the free store" by the **new** operator
 - The **new** operator returns a pointer to the allocated memory
 - A pointer is the address of the first byte of the memory
 - For example
 - **int* p = new int;** *// allocate one uninitialized **int***
 *// **int*** means "pointer to **int**"*
 - **int* q = new int[7];** *// allocate seven uninitialized **ints***
 *// "an array of 7 **ints**"*
 - **double* pd = new double[n];** *// allocate **n** uninitialized **doubles***
 - A pointer points to an object of its specified type
 - A pointer does **not** know how many elements it points to



Why use free store or heap?

- To allocate objects that have to outlive the function that creates them:
 - For example

```
double* make(int n)    // allocate n ints  
{  
    return new double[n];  
}
```

- Another example: vector's constructor
- Huge data (usually greater than >1MB)

Pointer Danger (when using new)



- Individual elements

```
int* p1 = new int;
```

// get (allocate) a new uninitialized int

```
int* p2 = new int(5);
```

// get a new int initialized to 5

```
int x = *p2;
```

// get/read the value pointed to by p2

// (or “get the contents of what p2 points to”)

// in this case, the integer 5

```
int y = *p1;
```

// undefined: y gets an undefined value; don't do that

Stack and Heap Memory

```
int foo() {  
    int* ptr; //<--nothing allocated yet (excluding the pointer itself, which is  
               // allocated here on the stack).  
  
    bool b = true; // Allocated on the stack.  
  
    if(b) {  
        //Create 500 bytes on the stack  
        int buffer[500];  
        //Create 500 bytes on the heap  
        ptr = new int[500];  
    } //<-- buffer is deallocated here, pBuffer is not  
  
    return 0;  
}  
//<--- oops there's a memory leak, I should have called  
// delete[] ptr;
```

A problem: memory leak

```
double* calc(int result_size, int max) // function returns pointer to a double
{
    double* p = new double[max];      // allocate another max doubles
                                     // i.e., get max doubles from the free store
    double* result = new double[result_size];
    // ... use p to calculate results to be put in result ...
    return result;
}
```

```
double* r = calc(200,100);           // oops! We "forgot" to give the memory
                                     // allocated for p back to the free store
```

- Lack of de-allocation (usually called "memory leaks") can be a serious problem in real-world programs
- A program that must run for a long time can't afford any memory leaks

A problem: memory leak

```
double* calc(int result_size, int max)
{
    int* p = new double[max];  // allocate another max doubles
                                // i.e., get max doubles from the free store
    double* result = new double[result_size];
    // ... use p to calculate results to be put in result ...
    delete[ ] p;                // de-allocate (free) that array
                                // i.e., give the array back to the free store

    return result;
}

double* r = calc(200,100);
// use r
delete[ ] r;                    // easy to forget
```

Memory leaks


- A program that needs to run "forever" can't afford any memory leaks
 - An operating system is an example of a program that “runs forever”
- If a function leaks 8 bytes every time it is called, how many days can it run before it has leaked/lost a megabyte?
 - Trick question: not enough data to answer, but about 130,000 calls
- All memory is returned to the system at the end of the program
 - If you run using an operating system (Windows, Unix, whatever)
- Program that runs to completion with predictable memory usage may leak without causing problems
 - *i.e.*, memory leaks aren't “good/bad” but they can be a major problem in specific circumstances

Memory leaks

- Another way to get a memory leak

```
void f()
{
    double* p = new double[27];
    // ...
    p = new double[42];
    // ...
    delete[] p;
}
```

// 1st array (of 27 doubles) leaked



Memory leaks

- How do we systematically and simply avoid memory leaks?
 - don't mess directly with **new** and **delete**
 - Use **vector**, etc.
 - Or use a garbage collector
 - A garbage collector is a program that keeps track of all of your allocations and returns unused free-store allocated memory to the free store (not covered in this course; see <http://www.research.att.com/~bs/C++.html>)
 - Unfortunately, even a garbage collector doesn't prevent all leaks
 - Use `smart_ptr` (`auto`, `unique_ptr`, `shared_ptr`)

Vector class: memory leak

```
class vector;
```

```
void f(int x)
```

```
{
```

```
    vector v(x);    // define a vector
```

```
    // (which allocates x doubles on the free store)
```

```
    // ... use v ...
```

```
    // give the memory allocated by v back to the free store
```

```
    // but how? (vector's elem data member is private)
```

```
}
```

Vector (destructor)

*// a very simplified **vector** of **doubles**:*

```
class vector {  
    int sz;                // the size  
    double* elem;          // a pointer to the elements  
public:  
    vector(int s)          // constructor: allocates/acquires memory  
        :sz(s), elem(new double[s]) { }  
    ~vector()              // destructor: de-allocates/releases memory  
        { delete[ ] elem; }  
    // ...  
};
```

- **Note:** this is an example of a general and important technique:
 - acquire resources in a constructor
 - release them in the destructor
- Examples of resources: memory, files, locks, threads, sockets

A problem: memory leak

```
void f(int x)
{
    int* p = new int[x]; // allocate x ints
    vector v(x);          // define a vector (which allocates another x ints)
    // ... use p and v ...
    delete[ ] p; // deallocate the array pointed to by p
    // the memory allocated by v is implicitly deleted here by vector's destructor
}
```

- The **delete** now looks verbose and ugly
 - How do we avoid forgetting to **delete[] p**?
 - Experience shows that we often forget
- Prefer **deletes** in destructors

Free store summary

- Allocate using **new**
 - New allocates an object on the free store, sometimes initializes it, and returns a pointer to it
 - `int* pi = new int;` *// default initialization (none for **int**)*
 - `char* pc = new char('a');` *// explicit initialization*
 - `double* pd = new double[10];` *// allocation of (uninitialized) array*
 - New throws a **bad_alloc** exception if it can't allocate
- Deallocate using **delete** and **delete[]**
 - **delete** and **delete[]** return the memory of an object allocated by **new** to the free store so that the free store can use it for new allocations
 - `delete pi;` *// deallocate an individual object*
 - `delete pc;` *// deallocate an individual object*
 - `delete[] pd;` *// deallocate an array*
 - Delete of a zero-valued pointer ("the null pointer") does nothing
 - `char* p = 0;`
 - `delete p;` *// harmless*

Warnings

- The description of memory layout is an *example*
 - It describes an *over-simplified* Unix layout, and current systems are ***much*** more complicated
 - Never write programs that assume any particular layout
 - Even the stack may grow both up and down!
- You can run out of a type of memory (stack, heap or other) long before you run out of memory
 - The simplest solution is to compile for 64-bit addressing
 - 32-bit desktop systems have been obsolete for over 5 years
 - That has ***NOTHING*** to do with how much memory they have
 - Generally, use a 64-bit system if you have 1 GB of memory or more

References

- Bjarne Stroustrup C++
- C++ Primer Plus
- Object Oriented Programming in C++
- <http://people.ds.cam.ac.uk/nmm1/C++/index.html>
(exercise 17)
- <http://www.thecodingforums.com/threads/stack-or-heap-for-c-object.689680/>
- <https://stackoverflow.com/questions/79923/what-and-where-are-the-stack-and-heap#79936>