C++ Parallelisation

Concurrency

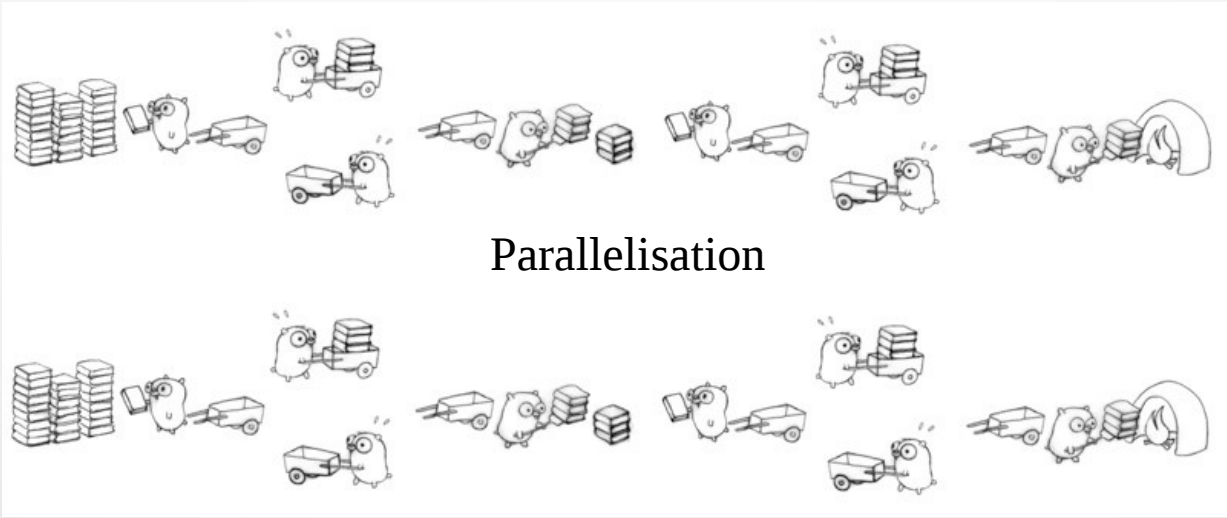

Parallelisation

Krishna Kumar

Shared memory

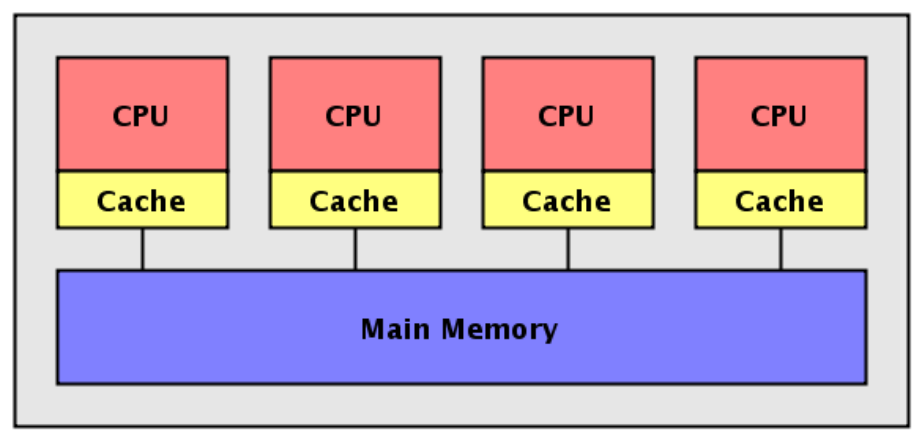

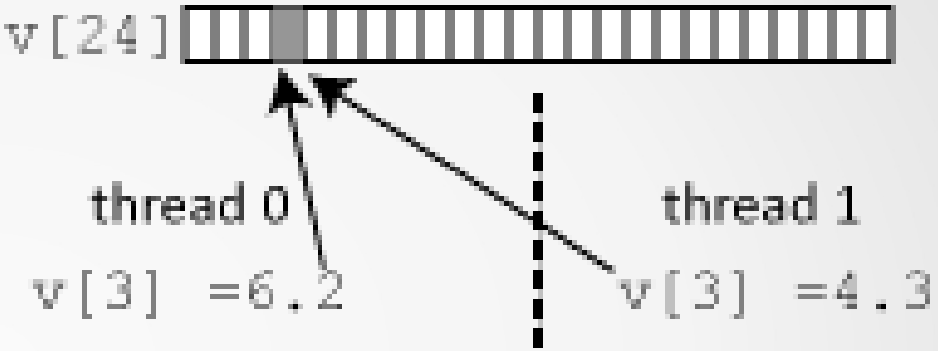

- two threads of execution can both address the same variables in a uniform manner, hereby assigning to an element of a vector whose components are in the virtual memory of the task.
- If the programmer wants thread 0 to use the value placed in the array by thread 1, he needs to use a mechanism which assures him that thread 1 has written the value before thread 0 reads it.

Distributed memory

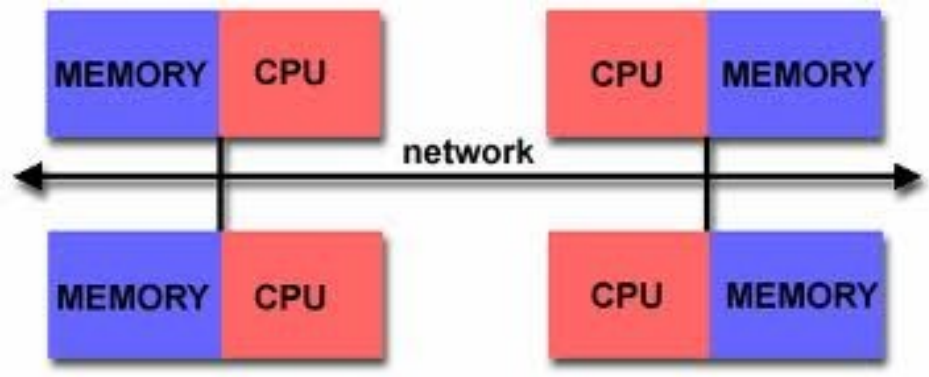

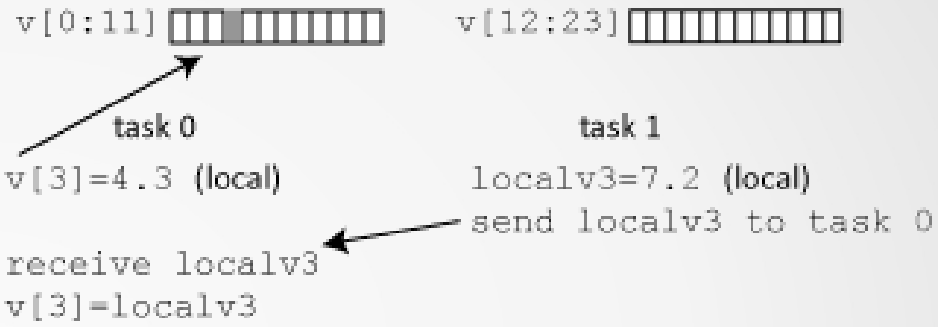

- Each task owns part of the data, and other tasks must send a message to the owner in order to update that data.
- These may be two tasks on the same computer so that they could just share memory, but the programmer is treating them as though they were not.
- Virtual address space is not shared.

C++ libraries for parallelisation

- **Shared memory**
  - OpenMP
  - C++11 Threads
  - Posix Threads
  - Intel TBB
- **Distributed Memory**
  - MPI
- **GPU**
  - CUDA
  - OpenCL
  - OpenACC
  - AMP

- Open Multi-Processing is an API to explicitly direct multi-threaded, shared memory parallelism
- Comprised of three primary API components:
  - Compiler Directives – Pragmas (pre-processor macros)
  - Runtime Library Routines
  - Environment Variables
- OpenMP is not:
  - For distributed memory parallel systems (by itself)
  - Guaranteed to make the most efficient use of shared memory
- **The programmer is responsible for synchronizing input and output.**

Thread based parallelism

- OpenMP programs accomplish parallelism exclusively through the use of threads.
- A thread of execution is the smallest unit of processing that can be scheduled by an operating system.
- Threads exist within the resources of a single process. Without the process, they cease to exist.
- Typically, the number of threads match the number of machine processors/cores. However, the actual use of threads is up to the application.

Fork – Join Model

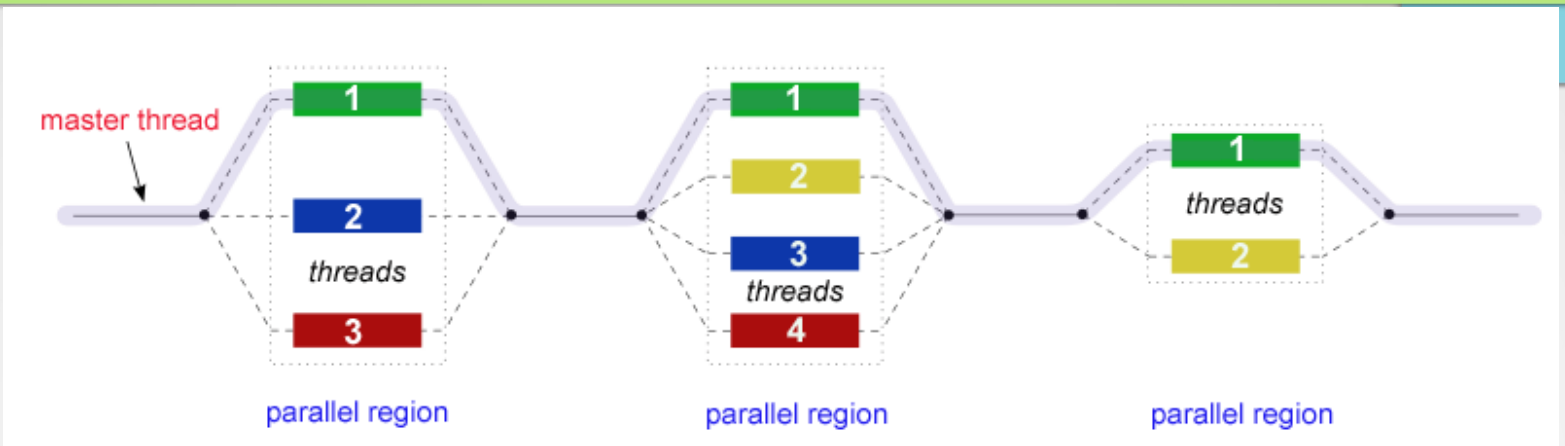

- FORK: the master thread then creates a team of parallel threads.
- The statements in the program that are enclosed by the parallel region construct are then executed in parallel among the various team threads.
- JOIN: When the team threads complete the statements in the parallel region construct, they synchronize and terminate, leaving only the master thread.
- The number of parallel regions and the threads that comprise them are arbitrary.

OpenMP: General code structure

```
#include <omp.h>
int main ()  {
      int var1, var2, var3;
      // Serial code . . .
      // Beginning of parallel section. Fork a team of threads.
      //Specify variable scoping
#pragma omp parallel private(var1, var2) shared(var3) {
      //  Parallel section executed by all threads
      // Run-time Library calls
      // All threads join master thread and disband
  }
// Resume serial code ..
}
```