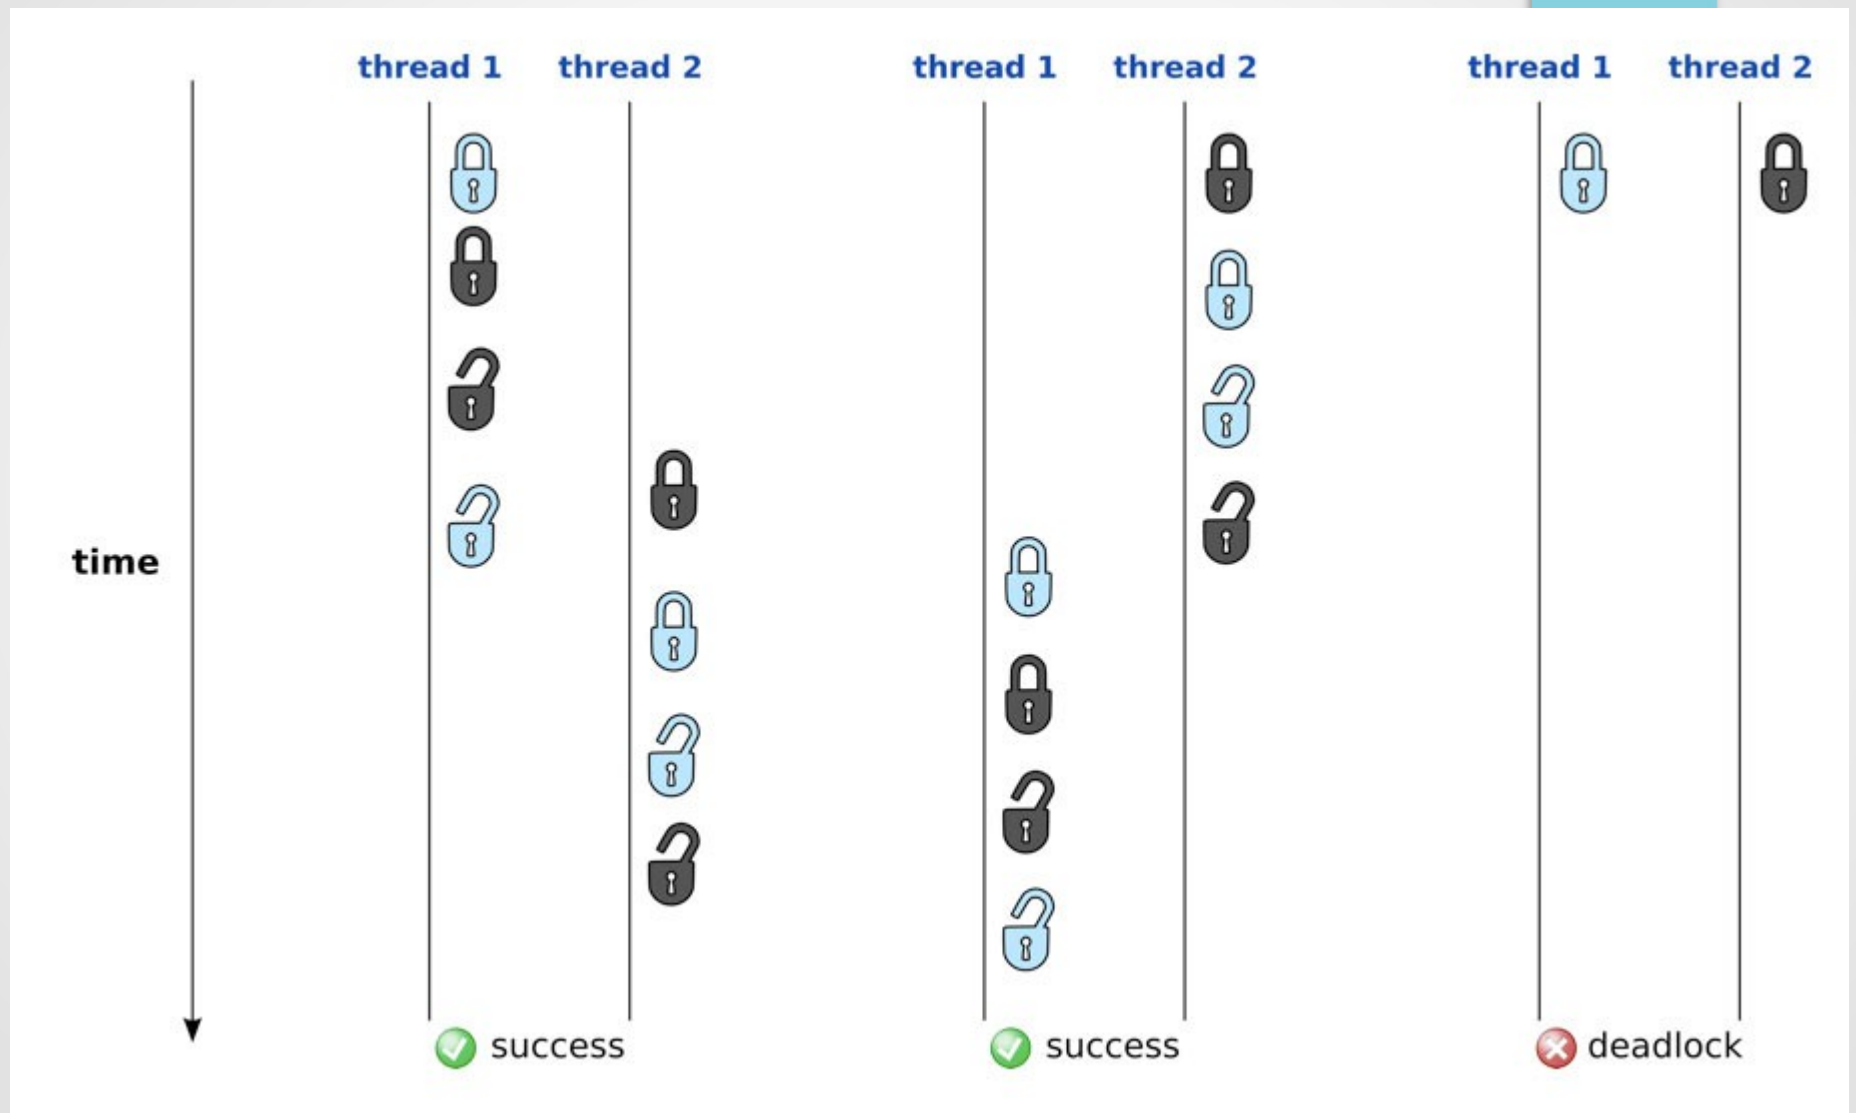


C++ Threads, Atomic & Mutex



Krishna Kumar

C++11 (Threads)

- Perhaps one of the biggest change to the language is the addition of multithreading support. Before C++11 – OpenMP, MPI to target multicore systems.
- **std::threads**
 - The class thread represents a single thread of execution. Threads allow multiple pieces of code to run asynchronously and simultaneously.
- Starting a new thread is very easy. When you create an instance of a `std::thread`, it will automatically be started.
- When you create a thread you have to give it the code it will execute. The first choice for that, is to pass it a function pointer.
- Calling `join()` function forces the current thread to wait for the other one (in this case, the main thread has to wait for the thread `t1` to finish). If you omit this call, the result is undefined.

Synchronisation

```
class Counter {  
    int value_;  
public:  
    Counter() : value_(0) {}  
    void increment() { ++value_; }  
    int value() { return value_; }  
    void value(int& newvalue)  
    { value_ = newvalue; }  
};
```

```
for (int i = 0; i < 500; ++i) {  
  
    threads.push_back(std::thread([  
        &counter]() {  
            for (int i = 0; i < 100; ++i)  
                counter.increment();  
        }));  
}  
  
// After thread.join()  
  
std::cout << counter.value() ;
```

Synchronisation issues

- The problem is that the incrementation is not an atomic operation. As a matter of fact, an incrementation is made of three operations:
 - Read the current value of value
 - Add one to the current value
 - Write that new value to value
- When you run that code using a single thread, there are no problems. But on several threads:
 - Thread 1 : read the value, get 0, add 1, so value = 1
 - Thread 2 : read the value, get 0, add 1, so value = 1
 - Thread 1 : write 1 to the field value and return 1
 - Thread 2 : write 1 to the field value and return 1
- These situations come from what we call interleaving.

Semaphores



Resources



Semaphore



Processes

C++ Mutex

- The C+11 standard provides `std::mutex` primitive. A mutex object also provides member functions – `lock()` and `unlock()` - to explicitly lock or unlock a mutex.

```
void increment() {  
    mutex.lock();  
    ++value;  
    mutex.unlock();  
}
```

- The first one enable a thread to obtain the lock and the second releases the lock. The `lock()` method is blocking. The thread will only return from the `lock()` method when the lock has been obtained.

Locks and Exceptions

- Imagine that the Counter has a decrement operation that throws an exception if the value is 0:

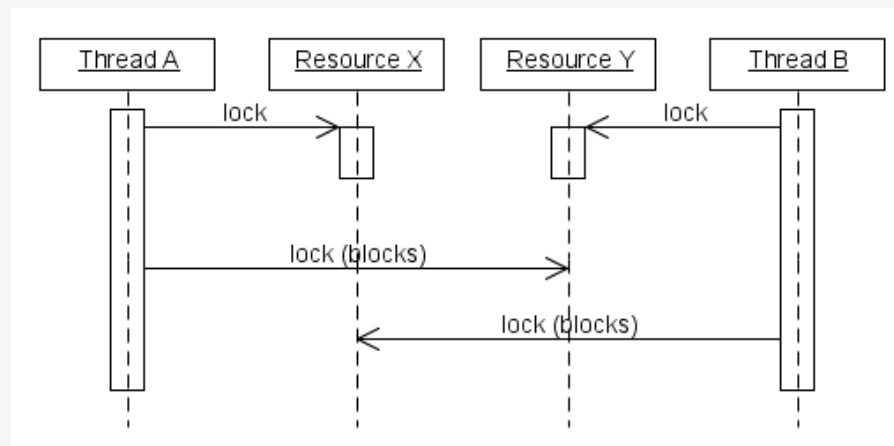
```
void decrement(){  
    if(value == 0){  
        throw "Value cannot be less than 0";  
    }  
    --value;  
}
```

- This wrapper works well in most of the cases, but when an exception occurs in the decrement method, you have a big problem. Indeed, if an exception occurs, the unlock() function is not called and so the lock is not released. As a consequence, your program is completely blocked.

Atomic

- The C++11 Concurrency Library introduces Atomic Types as a template class: `std::atomic`. You can use any Type you want with that template and the operations on that variable will be atomic and so thread-safe.
- If you want to make a big type (let's say 2MB storage), you can use `std::atomic` as well, but mutexes will be used. In this case, there will be no performance advantage.
- The main functions that `std::atomic` provide are the store and load functions that atomically set and get the contents of the `std::atomic`.

Recursive Lock – Deadlock!



References

- <http://baptiste-wicht.com/posts/2012/03/cpp11-concurrency-part1-start-threads.html>
- <http://baptiste-wicht.com/posts/2012/03/cpp11-concurrency-tutorial-part-2-protect-shared-data.html>
- <http://baptiste-wicht.com/posts/2012/04/cpp11-concurrency-tutorial-advanced-locking-and-condition-variables.html>