

Name Lookup and the Interface Principle

```
namespace A {  
    struct X;  
    void g( X );  
}  
namespace B {  
    void g( A::X x ) {  
        g( x ); // which g()?  
    }  
}
```



Krishna Kumar

Name Lookup

- When you call a function, which function do you call? The answer is determined by name lookup, but you're almost certain to find some of the details surprising.
 - In the following code, which functions are called? Why?
- ```
namespace A {
 – struct X;
 – struct Y;
 – void f(int);
 – void g(X);
}

namespace B {
 – void f(int i) { f(i); } // which f()?
 – void g(A::X x) { g(x); }
 // which g()?
 – void h(A::Y y) { h(y); }
 // which h()?
}
```

# Name lookup

- **Let's start simple.**
- namespace A {
  - struct X;
  - void f( int );
- }
- namespace B {
  - void f( int i ) { f( i ); } // which f()?
- }
- This f() calls itself, with infinite recursion. The reason is that the only visible f() is B::f() itself.
- There is another function with signature f(int) , namely the one in namespace A . If B had written "using namespace A; " or " using A::f; ", then A::f(int) would have been visible as a candidate when looking up f(int) , and the f(i) call would have been ambiguous between A::f(int) and B::f(int) . Since B did not bring A::f(int) into scope, however, only B::f(int) can be considered, so the call unambiguously resolves to B::f(int) .

# Koenig Lookup

- A good knowledge of C++'s name lookup rules—in particular, **Koenig lookup (argument-dependent lookup – ADL)** is required.
- ```
void g( A::X x ) {  
    – g( x ); // which g()?  
}
```
- This call is ambiguous between `A::g(X)` and `B::g(X)` . The programmer must explicitly qualify the call with the appropriate namespace name to get the `g()` they want.
- You may at first wonder why this should be so. After all, as with `f()` , **B** hasn't written " using " anywhere to bring `A::g(X)` into its scope, so you'd think that only `B::g(X)` would be visible, right? Well, this would be true except for an extra rule that C++ uses when looking up names.

Koenig Lookup (cont...)

- Koenig Lookup (simplified):
 - *“If you supply a function argument of class type (here x , of type $A::X$), then to look up the correct function name the compiler considers matching names in the namespace (here A) containing the argument's type.”*
- Here is an example:
- namespace std {
 - class string { };
 - operator<<(T);
- }
- int main() {
 - std::string parm = “Hello world!”;
 - std::cout << parm ; // OK, calls std::string's <<
- }

Name lookup (cont)

- Now back to a simple example:
- `void h(A::Y y) {`
 - `h(y); // which h()?`
- `}`
- There is no other `h(A::Y)` , so this `h()` calls itself with infinite recursion.
- Although `B::h()` 's signature mentions a type found in namespace `A` , this doesn't affect name lookup because there are no functions in `A` matching the name and signature `h(A::Y)` .
- So, what does it mean? In short, the meaning of code in namespace `B` is being affected by a function declared in the completely separate namespace `A` , even though `B` has done nothing but simply mention a type found in `A` and there's no "using" in sight.
- What this means is that namespaces aren't quite as independent as people originally thought. Don't stardenouncing namespaces just yet, though; namespaces are still pretty independent, and they fit their intended uses to a T.

Class interface

- A traditional definition of a class:
 - A class describes a set of data, along with the functions that operate on that data.
- Question: What functions are "part of " a class, or make up the interface of a class?
 - Hint #1: Clearly nonstatic member functions are tightly coupled to the class, and public nonstatic member functions make up part of the class's interface. What about static member functions? What about free functions?
 - Hint #2: consider the implications of name lookup.

Class interface

A class describes a set of data, along with the functions that operate on that data.

- Programmers often unconsciously misinterpret this definition, saying instead: "Oh yeah, a class, that's what appears in the class definition—the member data and the member functions." But that's not the same thing, because it limits the word functions to mean just member functions. Consider:
- ***class X { /*...*/ };***
- ***void f(const X&);***
- The question is: Is **f** part of **X** ? Some people will automatically say "No" because **f** is a nonmember function (or free function). Others might realize something fundamentally important: If the code appears together in one header file, it is not significantly different from:
- ***class X {***
 - ***/*...*/***
 - ***public:***
 - void f() const;***
- ***};***
- Besides access rights, **f** is still the same, taking a pointer/reference to **X** . The *this* parameter is just implicit in the second version, that's all.

Class interface principle

- So, if example code all appears in the same header, we're already starting to see that even though **f** is not a member of **X**, it's nonetheless strongly related to **X**.
- On the other hand, if **X** and **f** do not appear together in the same header file, then **f** is just some old client function, not a part of **X** (even if **f** is intended to augment **X**). We routinely write functions with parameters whose types come from library headers, and clearly our custom functions aren't part of those library classes.
- With that example in mind, Interface Principle is:
 - For a class **X**, all functions, including free functions, that both
 - "Mention" **X**
 - Are "supplied with" **X**
 - are logically part of **X**, because they form part of the interface of **X**.
- By definition, every member function is "part of" **X** :
 - Every member function must "mention" **X** - a nonstatic member function has an implicit **this** parameter of type **X* const** or **const X* const** ; a static member function is in the scope of **X**
 - Every member function must be "supplied with" **X** (in **X** 's definition).

Class interface principle (cont...)

- Applying the Interface Principle to the example code gives the same result as our original analysis. Clearly, **f** mentions **X**. If **f** is also "supplied with" **X** (same header file and/or namespace), then according to the Interface Principle, **f** is logically part of **X** because it forms part of the interface of **X**.
- Interface Principle behaves in exactly the same way that Koenig lookup does.
- So the Interface Principle is a useful touchstone to determine what is really "part of" a class. Is it unintuitive that a free function should be considered part of a class? Then let's give real weight to this example by giving a more common name to **f**.
- `class X { /*...*/ };`
- `/*...*/`
- `ostream& operator<<(ostream&, const X&);`
- Here the Interface Principle's rationale is perfectly clear, because we understand how this particular freefunction works. If `operator<<` is "supplied with" **X** (for example, in the same header and/or namespace), then **operator<<** is logically part of **X** because it forms part of the interface of **X**. That makes sense even though the function is a nonmember, because we know that it's common practice for a class's author to provide **operator<<**. If, instead, **operator<<** comes, not from **X**'s author, but from client code, then it's not part of **X** because it's not "supplied with" **X**.

OO Interface principle

- // Example 1

- struct _iobuf { /*...data goes here...*/ };
- typedef struct _iobuf FILE;
- FILE* fopen (const char* filename, const char* mode);
- Int fclose(FILE* stream);
- Int fseek (FILE* stream,
 - long offset, int origin);
- long ftell (FILE* stream);
- /* etc. */

- // Example 2

- class FILE {
- public:
 - FILE(const char* filename, const char* mode);
 - ~FILE();
 - int fseek(long offset, int origin);
 - long ftell();
- private:
 - /*...data goes here...*/
- };

FILE* parameters have just become implicit this parameters. Here it's clear that fseek is part of FILE , just as it was in Example 1, even though there it was a nonmember. We can even merrily make some functions members and some not.

More Koenig Lookup

- namespace NS {
 - class T { }; // some header
 - }
 - void f(NS::T);
 - int main() {
 - NS::T parm;
 - f(parm); // OK: calls global f
 - }
- namespace NS {
 - class T { };
 - void f(T); // <-- new function
 - }
 - void f(NS::T);
 - int main() {
 - NS::T parm;
 - f(parm);
 - // ambiguous: NS::f or Global f?
 - }

Adding a function in a namespace scope "broke" code outside the namespace, even though the client code didn't write using to bring NS 's names into its scope!

More Koenig lookup (Myer's example)

- Namespace A {
 - class X {};
- }
- namespace B {
 - void f(A::X);
 - void g(A::X parm) {
 - f(parm);
 - // OK: calls B::f
 - }
- }

- namespace A {
 - class X { };
 - void f(X); // <-- new function
- }
- namespace B {
 - void f(A::X);
 - void g(A::X parm) {
 - f(parm);
 - // ambiguous: A::f or B::f?
 - }
- }

More Koenig loopup (Myer's)

- "The whole point of namespaces is to prevent name collisions, isn't it? But adding a function in one namespace actually seems to 'break' code in a completely separate namespace."
- True, namespace **B** 's code seems to break merely because it mentions a type from **A** . **B** 's code didn't write a using namespace **A**; anywhere. It didn't even write using **A::X**;
- This is not a problem, and **B** is not broken. This is in fact exactly what should happen. If there's a function **f(X)** in the same namespace as **X** , then, according to the Interface Principle, **f** is part of the interface of **X** .
- It doesn't matter that **f** happens to be a free function; to see clearly that it's nonetheless logically part of **X** , just give it another name.

Interface principle

- **In general, if A and B are classes and A::g(B) is a member function of A :**
- Because A::g(B) exists, clearly A always depends on B . No surprises so far.
- If A and B are supplied together, then of course A::g(B) and B are supplied together. Therefore, because A::g(B) both "mentions" B and is "supplied with" B , then according to the Interface Principle, it follows that A::g(B) is part of B , and because A::g(B) uses an (implicit) A* parameter, B depends on A . Because A also depends on B , this means that A and B are interdependent.
- At first, it might seem like a stretch to consider a member function of one class as also part of another class, but this is true only if A and B are also supplied together.
- **Unlike classes, namespaces don't need to be declared all at once, and what's "supplied together" depends on what parts of the namespace are visible.**
- `///---file a.h---`
- `namespace N { class B; }// forward decl`
- `namespace N { class A; }// forward decl`
- `class N::A { public: void g(B); };` In a.h: A and B are supplied together and are interdependent.
- `///---file b.h---`
- `namespace N { class B { /*...*/ }; }` Clients of B include b.h , A and B are not supplied together.

Name Hiding

- `// from a base class`
- `struct B {`
 - `int f(int);`
 - `int f(double);`
 - `int g(int);`
- `};`
- `struct D : public B {`
 - `private:`
 - `int g(std::string, bool);`
- `};`
- `D d;`
- `int i;`
- `d.f(i); // ok, means B::f(int)`
- `d.g(i); // error: g takes 2 args`
- In short, when we declare a function named **g** in the derived **class D**, it automatically hides all functions with the same name in all direct and indirect base classes. It
- doesn't matter a whit that **D::g** "obviously" can't be the function that the programmer meant to call (not only does **D::g** have the wrong signature, but it's private and, therefore, inaccessible to boot), because **B::g** is hidden and can't be considered by name lookup.

Name Hiding

- To see what's really going on, let's look in a little more detail at what the compiler does when it encounters the function call **d.g(i)**.
- First, it looks in the immediate scope, in this case the scope of class **D**, and makes a list of all functions it can find that are named **g** (regardless of whether they're accessible or even take the right number of parameters). Only if it doesn't find any at all does it then continue "outward" into the next enclosing scope and repeat—in this case, the scope of the base class **B**—until it eventually either runs out of scopes without having found a function with the right name or else finds a scope that contains at least one candidate function.
- If a scope is found that has one or more candidate functions, the compiler then stops searching and works with the candidates that it's found, performing overload resolution and then applying access rules.
- It makes intuitive sense that a member function that's a near-exact match ought to be preferred over a global function that would have been a perfect match had we considered the parameter types only.

Working around unwanted Name Hiding

- **// Asking for a name**
- **// from a base class**
- **D d;**
- **int i;**
- **d.f(i);** // ok, means **B::f(int)**
- **d.B::g(i);**
- **// ok, asks for B::g(int)**
- First, the calling code can simply say which one it wants and force the compiler to look in the right scope.
- **// Un-hiding a name**
- **// from a base class**
- **struct D : public B {**
 - **using B::g;**
- **private:**
 - **int g(std::string, bool);**
- **};**
- usually more appropriate, the designer of class D can make **B::g** visible with a using declaration. This allows the compiler to consider **B::g** in the same scope as **D::g** for the purposes of name lookup and subsequent overload resolution.

Namespaces and the Interface Principle

- **// Will this compile?**
- // In some library header:
 - namespace N { class C {};
 - int operator+(int i, N::C) {
 - return i+1;
 - }
- // A mainline to exercise it:
 - #include <numeric>
 - int main() {
 - N::C a[10];
 - std::accumulate(a, a+10, 0);
 - }
- **The std::accumulate template looks something like this:**
- namespace std {
 - template<class Iter, class T>
 - inline T accumulate(Iter first, Iter last, T value) {
 - while(first != last) {
 - value = value + *first; // 1
 - ++first;
 - }
 - return value;
 - }

Namespaces and the Interface Principle

- The code actually calls `std::accumulate<N::C*,int>` . In line 1 above, how should the compiler interpret the expression value `+ *first` ? Well, it's got to look for an `operator+()` that takes an `int` and an `N::C` (or parameters that can be converted to `int` and `N::C`).
- It just so happens that we have just such an `operator+(int,N::C)` at global scope! Look, there it is! Cool. So everything must be fine, right?
- The problem is that the compiler may or may not be able to see the `operator+(int,N::C)` at global scope, depending on what other functions have already been seen to be declared in namespace `std` at the point where `std::accumulate<N::C*,int>` is instantiated.
- The compilation depends entirely on whether this implementation's version of the standard header `numeric` : a) declares an `operator+()` (any `operator+()` , suitable or not, accessible or not); or b) includes any other standard header that does so. Unlike standard C, standard C++ does not specify which standard headers will include each other, so when you include `numeric`, you may or may not get header `iterator` too which does define several `operator+()` functions.

Some Fun with Compilers

- It's bad enough that the compiler can't find the right function if there happens to be another `operator+()` in the way, but typically the `operator+()` that does get encountered in a standard header is a template, and compilers generate notoriously difficult-to-read error messages when templates are involved.
 - **error C2784:** 'class `std::reverse_iterator`<'template-parameter-1','template-parameter-2','template-parameter-3','template-parameter-4','template-parameter-5'> `__cdecl std::operator+(template-parameter-5,const class std::reverse_iterator`<'template-parameter-1','template-parameter-2','template-parameter-3','template-parameter-4','template-parameter-5'>&)' : could not deduce template argument for 'template-parameter-5' from 'int'. -
 - The compiler is merely complaining (as clearly as it can) that it did find an `operator+()` but can't figure out how to use it in an appropriate way.
 - **error C2677:** binary '+' : no global operator defined which takes type 'class `N::C`' (or there is no acceptable conversion)

The solution – Namespace & Interface

- When we encountered this problem in the familiar guise of base/derived name hiding, we had two possible solutions: Have the calling code explicitly say which function it wants, or write a using declaration to make the desired function visible in the right scope.
- Neither solution works in this case. The first is possible but places an unacceptable burden on the programmer; the second is impossible.
- First case requires the programmer to use the version of `std::accumulate` that takes a predicate and explicitly say which one he wants each time!

The solution - Namespace

- The real solution is to put our operator+() where it has always truly belonged and where it should have been put in the first place: in namespace N .
- // in some library header
- **namespace N{**
 - **class C {};**
 - **int operator+(int i, N::C) { return i+1; }**
- **}**
- **// a mainline to exercise it**
- **#include <numeric>**
- **int main() {**
 - **N::C a[10];**
 - **std::accumulate(a, a+10, 0); // now ok**
- **}**
- Now that the operator+() is in the same namespace as the second parameter, when the compiler tries to resolve the " + " call inside std::accumulate , it is able to see the right operator+() because of Koenig lookup.
- in addition to looking in all the usual scopes, the compiler shall also look in the scopes of the function's parameter types to see if it can find a match. N::C is in namespace N , so the compiler looks in namespace N , and happily finds exactly what it needs, no matter how many other operator+() 's happen to be lying around and cluttering up namespace std .

Conclusions

- The problem arose because the example did not follow the Interface Principle:
 - For a class X, all functions, including free functions, that both
 - "Mention" X
 - Are "supplied with" X
 - are logically part of X, because they form part of the interface of X.
- If an operation, even a free function (and especially an operator) mentions a class and is intended to form part of the interface of a class, then always be sure to supply it with the class—which means, among other things, to put it in the same namespace as the class.
- Use namespaces wisely. Either put all of a class inside the same namespace—including things that to innocent eyes don't look like they're part of the class, such as free functions that mention the class (don't forget the Interface Principle)—or don't put the class in a namespace at all.

References

- Exceptional C++ - Herb Sutter
-